

RAGHURAM P

Hyderabad, India
raghu.builds@gmail.com ◊ LinkedIn ◊ GitHub

Professional Summary

Full-Stack Engineer specializing in AI-native enterprise platforms, owning products end-to-end from Angular/RxJS frontends to async FastAPI services and event-driven real-time systems. Architected the SSE + Redis Streams backbone powering live code, wireframe, and app generation, and shipped multi-agent LLM products on CrewAI, AWS Bedrock, and Azure OpenAI across 5+ enterprise client environments.

Work Experience

Full-Stack Engineer — GenAI Platforms Ascendion	Jun 2024 – Present Hyderabad, India
- Co-built and production-hardened an internal multi-agent process-orchestration engine running multi-agent LLM workflows through human-in-the-loop approval gates, real-time streaming, and governed LLM routing across cloud providers — owning multi-pod streaming transport, LLM retry/resilience, conversation compaction, and per-run cost tracking; adopted by 2K+ users daily. - Co-built the backend for an internal AI coding assistant (VS Code plugin), re-architecting a single-agent prototype into a multi-agent orchestration system (Main-Agent + sub-agents on CrewAI flows, 30+ tools, skills & command frameworks) — adopted by 3K+ users daily across 5+ client environments. - Designed the SSE + Redis Streams real-time architecture that is the backbone of an AI app-generation platform, replacing polling with a decoupled pub/sub transport — per-session connection deduplication, exponential backoff, Last-Event-ID checkpointing, history backfill on zero-replay reconnect, and multi-pod fan-out. - Re-architected the entire frontend of an AI app-generation platform from React to Angular — standalone components, RxJS/Signals state orchestration, lazy-loaded feature modules, and shared stream-handling utilities — supporting 5+ generation workflows and cutting cross-team code duplication. - Co-built a prompt-to-React code-generation system converting wireframes into production-ready modular components and routing logic deployed to Azure — component synthesis, routing generation, and sandboxed preview, deployed to Azure. - Built async FastAPI services powering an AI product-planning platform’s multi-step flow (Brainstormer → Persona → SWOT → Roadmap → Summary) on PostgreSQL (asyncpg, AWS RDS IAM auth) with versioned SQL migrations and Redis pub/sub — delivering zero-touch Azure deployments. - Co-built a GenAI wireframe generator converting product briefs (PRDs, sketches, prompts) into production-ready UI artifacts — brief-to-artifact pipeline with structured extraction and streamed rendering; serving 500+ users daily. - Served as deployment point-of-contact across 5 enterprise client environments (hardware, banking, and fintech domains), resolving environment-specific deployment blockers and ensuring stable production go-lives. - Co-built a prompt-driven execution engine on LLM agent orchestration with RAG and ReAct pipelines, compressing workflow planning from 1.5 hours to 15 minutes for 1.5K+ users.	

Technical Skills

Programming Languages	Python, TypeScript, JavaScript, SQL
Frontend	Angular (Standalone Components, RxJS, Signals), React, Monaco Editor, SCSS, Tailwind CSS
Backend & APIs	FastAPI, gRPC, REST, SSE / Streaming APIs, Async Processing, Pydantic
AI / GenAI	Multi-Agent Orchestration (CrewAI), RAG, LLM Routing, AWS Bedrock, Azure OpenAI
Databases & Storage	PostgreSQL, Redis, MongoDB, SQLite (sqlite-vec), asyncpg, SQLAlchemy
Messaging & Real-Time	Redis Streams (Pub/Sub), SSE, WebSockets, Apache Kafka, Event-Driven Architecture
Cloud & DevOps	AWS, Azure (DevOps Pipelines, App Services, Artifacts), Docker, Kubernetes, CI/CD
Architecture & Testing	Distributed Systems, Microservices, Micro-Frontend, OpenTelemetry, pytest, Jest

Education

B.Tech, Information Technology , Vel Tech Multi Tech Engineering College CGPA: 7.78	2019 – 2023
---	-------------