

RAGHURAM P

Hyderabad, India

raghu.builds@gmail.com ♦ LinkedIn ♦ GitHub ♦ Portfolio

Professional Summary

Full-Stack Engineer specializing in AI-native enterprise platforms, owning products end-to-end from Angular/RxJS frontends to async FastAPI services and event-driven real-time systems. Architected the SSE + Redis Streams backbone powering live code, wireframe, and app generation, and shipped multi-agent LLM products on CrewAI, AWS Bedrock, and Azure OpenAI across 5+ enterprise client environments.

Work Experience

Full-Stack Engineer — GenAI Platforms

Jun 2024 – Present

Ascendion

Hyderabad, India

- Co-built and production-hardened an **internal multi-agent process-orchestration engine** running multi-agent LLM workflows through human-in-the-loop approval gates, real-time streaming, and governed LLM routing across cloud providers — owning multi-pod streaming transport, LLM retry/resilience, conversation compaction, and per-run cost tracking; adopted by **2K+ users daily**.
- Co-built the backend for an **internal AI coding assistant** (VS Code plugin), re-architecting a single-agent prototype into a **multi-agent orchestration system** (Main-Agent + sub-agents on CrewAI flows, 30+ tools, skills & command frameworks) — adopted by **3K+ users daily across 5+ client environments**.
- Designed the **SSE + Redis Streams** real-time architecture that is the backbone of an AI app-generation platform, replacing polling with a decoupled pub/sub transport — per-session connection deduplication, exponential backoff, **Last-Event-ID** checkpointing, history backfill on zero-replay reconnect, and multi-pod fan-out.
- Re-architected the entire frontend of an AI app-generation platform from **React to Angular** — standalone components, RxJS/Signals state orchestration, lazy-loaded feature modules, and shared stream-handling utilities — supporting 5+ generation workflows and cutting cross-team code duplication.
- Built async **FastAPI** services powering an AI product-planning platform's multi-step flow (Brainstormer → Persona → SWOT → Roadmap → Summary) on **PostgreSQL (asyncpg, AWS RDS IAM auth)** with versioned SQL migrations and Redis pub/sub — delivering zero-touch Azure deployments.
- Served as deployment point-of-contact across **5 enterprise client environments** (hardware, banking, and fintech domains), resolving environment-specific deployment blockers and ensuring stable production go-lives.

Open Source

tessera / verity / sentinel — correctness libraries for LLM systems: tool-call transcript repair, retrieval measurement, and human-in-the-loop approval gating. Property-tested with Hypothesis; 97–99% coverage against enforced floors; mypy strict; zero runtime dependencies. *Python*

atlas — raghuram.vercel.app: one typed content layer feeds every page, sitemap and `llms.txt`; the build fails on a dead link. Lighthouse 1.00 on performance, accessibility, best practices and SEO, enforced in CI. *Next.js, TypeScript*

Technical Skills

Programming Languages	Python, TypeScript, JavaScript, SQL
Frontend	Angular (Standalone Components, RxJS, Signals), React, Monaco Editor, SCSS, Tailwind CSS
Backend & APIs	FastAPI, gRPC, REST, SSE / Streaming APIs, Async Processing, Pydantic
AI / GenAI	Multi-Agent Orchestration (CrewAI), RAG, LLM Routing, AWS Bedrock, Azure OpenAI
Databases & Storage	PostgreSQL, Redis, MongoDB, SQLite (sqlite-vec), asyncpg, SQLAlchemy
Messaging & Real-Time	Redis Streams (Pub/Sub), SSE, WebSockets, Apache Kafka, Event-Driven Architecture
Cloud & DevOps	AWS, Azure (DevOps Pipelines, App Services, Artifacts), Docker, Kubernetes, CI/CD
Architecture & Testing	Distributed Systems, Microservices, Micro-Frontend, OpenTelemetry, pytest, Jest

Education

B.Tech, Information Technology, Vel Tech Multi Tech Engineering College

2019 – 2023

CGPA: 7.78